

# magnetic tape data recovery in software

a very nerdy talk for the Vintage Computer Festival West

Len Shustek

CHM Chairman emeritus

```

*****
*                                     *
*                                     *
*          PLAGO/360                  *
*                                     *
*          THE POLYTECHNIC LOAD-AND-GO PL/I TRANSLATOR *
*                                     *
*          VERSION I                  *
*                                     *
*****

```

The Polytechnic Institute of Brooklyn  
September 25, 1968

# the failure




Qty	Description	Unit Price	Line Total
1	Restore 9-track 800 BPI tape and land data to new media for Client delivery. This is a sole source of media being sent to eMag for processing; no other copies exist.	\$1,000 (flat fee)	\$1,000


```
*      <?PLAGO SYSTEM DUMP ROUTINE
*
*      ???THIS CSECT IS ENTERED BY THE E-MACHINE FOR THE INTERPRETATION
*      OF THE BUILT-IN-FUNCTION 'DEBUG'.
*
*      ???THIS FUNCTION IS INTENDED FOR USE ??LY BY THOSE KNOWLEDGEABLE
*      IN THE INTERNALS OF PLAGO.
*
*INPUT-  R4  OUT?UT BUFFER ADDRESS
*        R5  OUTPUT ROUTINE ADDRESS
*        ?  R15 DUMP ROUTINE BASE ADDRESS
*        ?  R6-R12 AS USED IN E-MACHINE
*
*      ?? IN ADDITION, THE TOP ELEMENT OF THE DSTK IS ASSUMED TO BE
*      ???POINTING AT A FIXED BINARY NUMBER WHOSE BITS, STARTING FROM
*      ??THE RIGHT, HAVE THE FOLLOWING SIGNIFICANCE:
*      ???1. REGISTER DISPLAY REQUESTED
*      ??2. D-STK DUMP REQUESTED
```

? marks errors

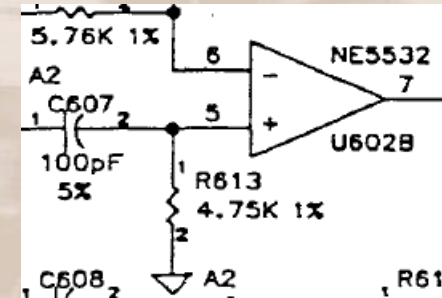
# traditional hardware




pre-amp




differential amp




differentiator




comparator



threshold detector



zero-crossing  
detector




data bit  
decoder

...and then block and error handling!



# software-based\*




pre-amp

analog to  
digital  
conversion

detect  
flux  
transitions

detect  
bits

assemble  
into words,  
then blocks



in software,  
so we can retry  
without rereading  
the tape!


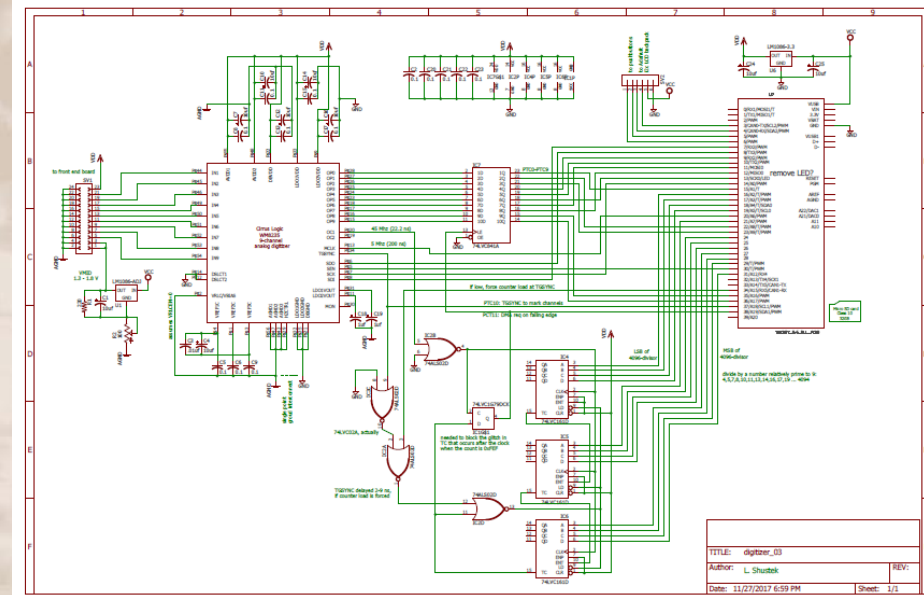
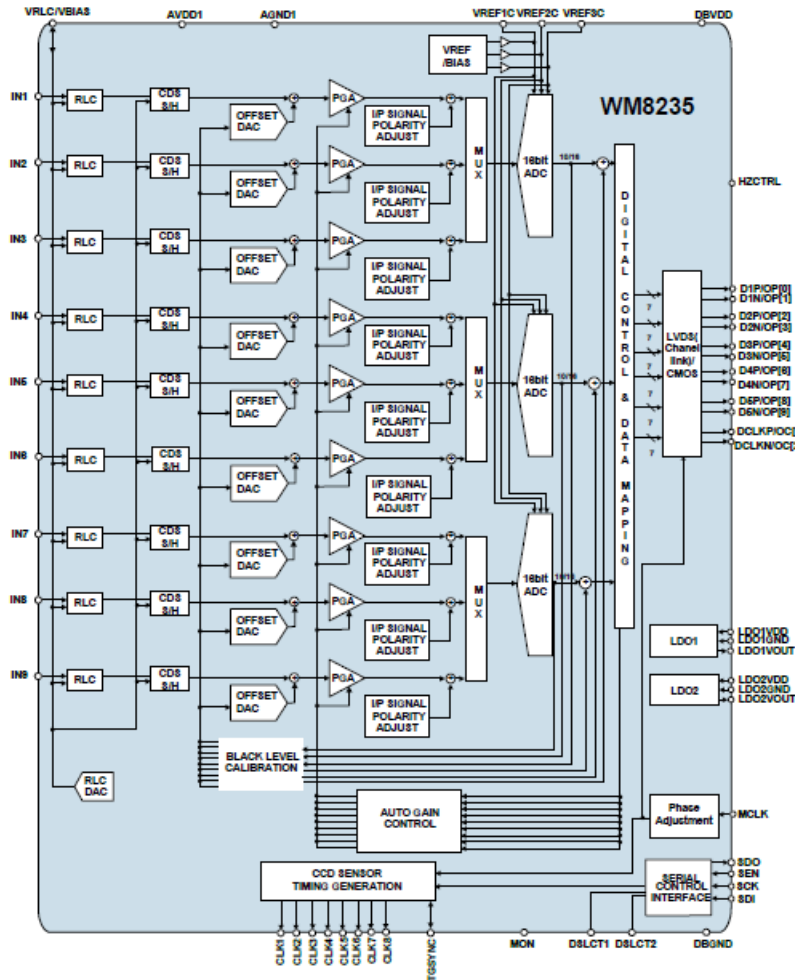
\* Inspired by Paul Pierce's  
Java-based project in 2004

# my first attempt




WM8235

## BLOCK DIAGRAM





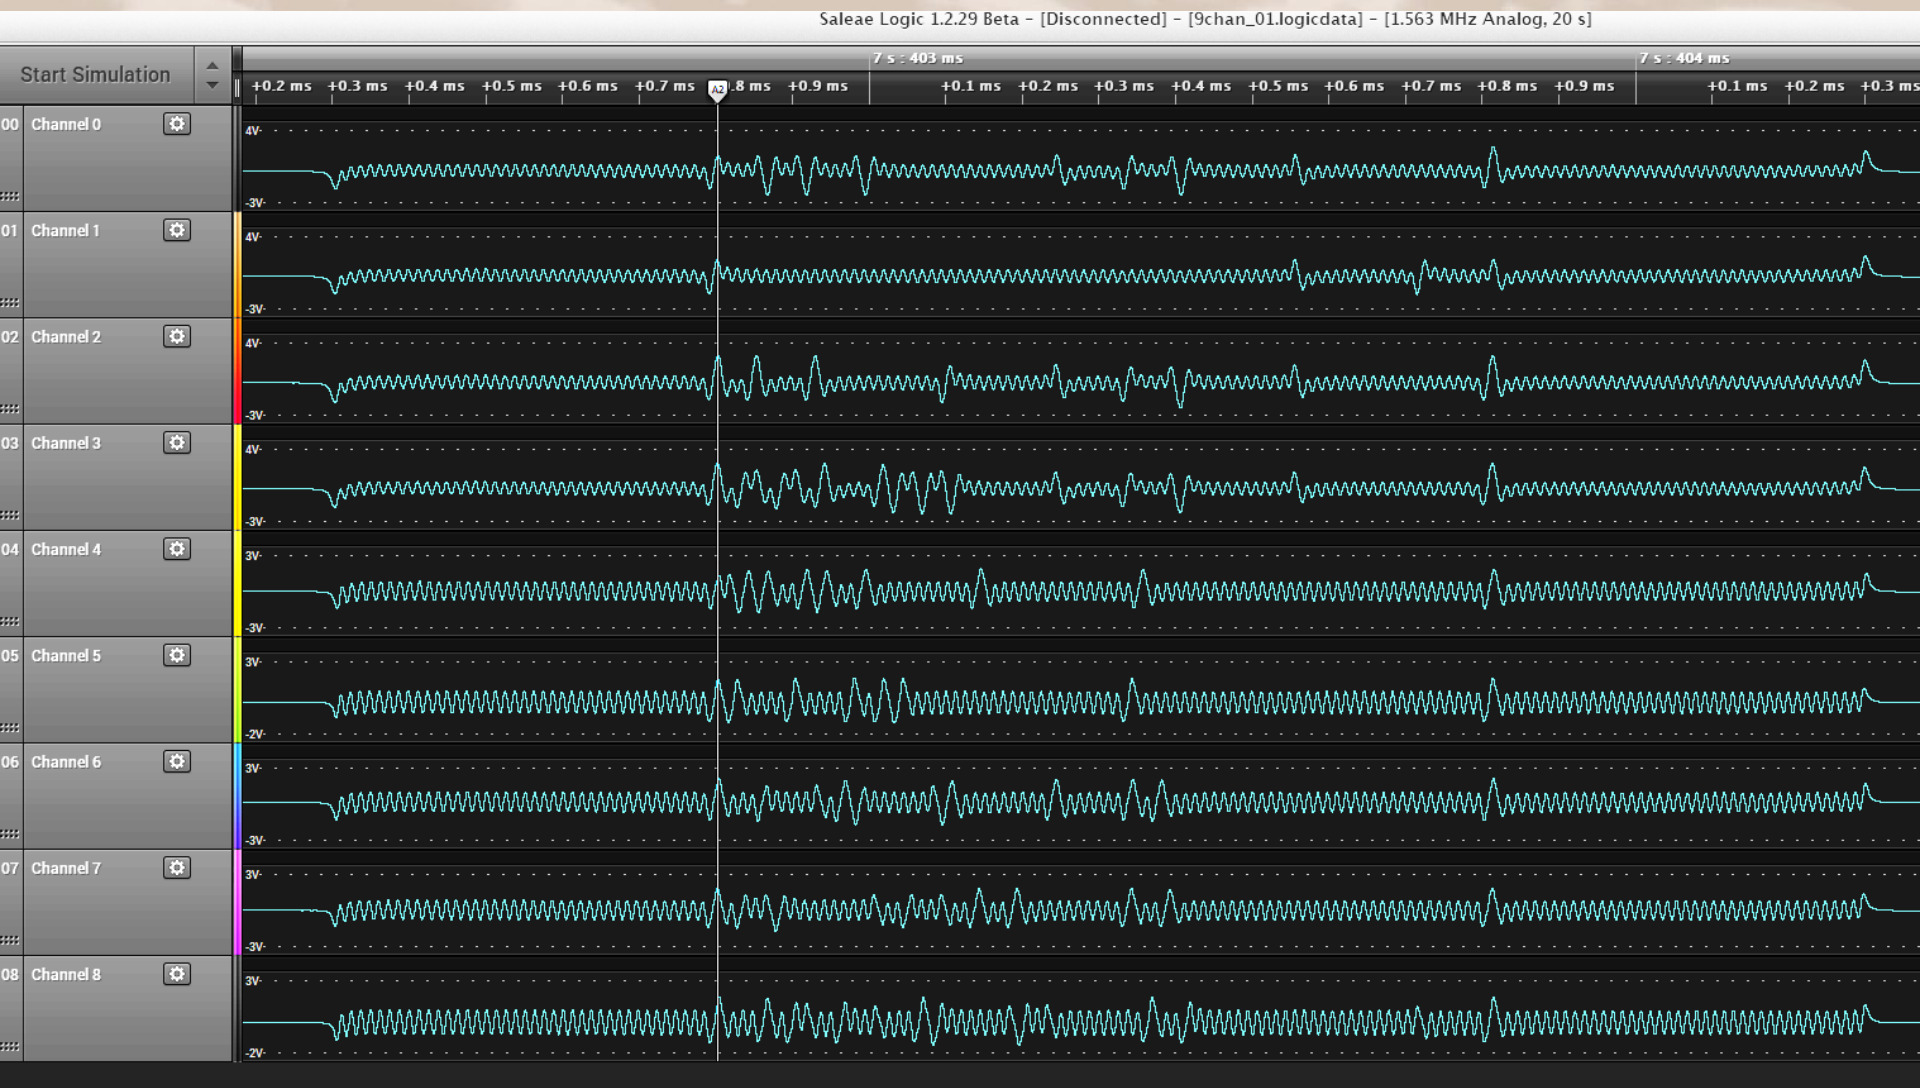


# easier:

## Saleae Logic Pro 16 logic analyzer

- *16 analog channels*
- *max 50M samples/sec*
- *“unlimited” storage*

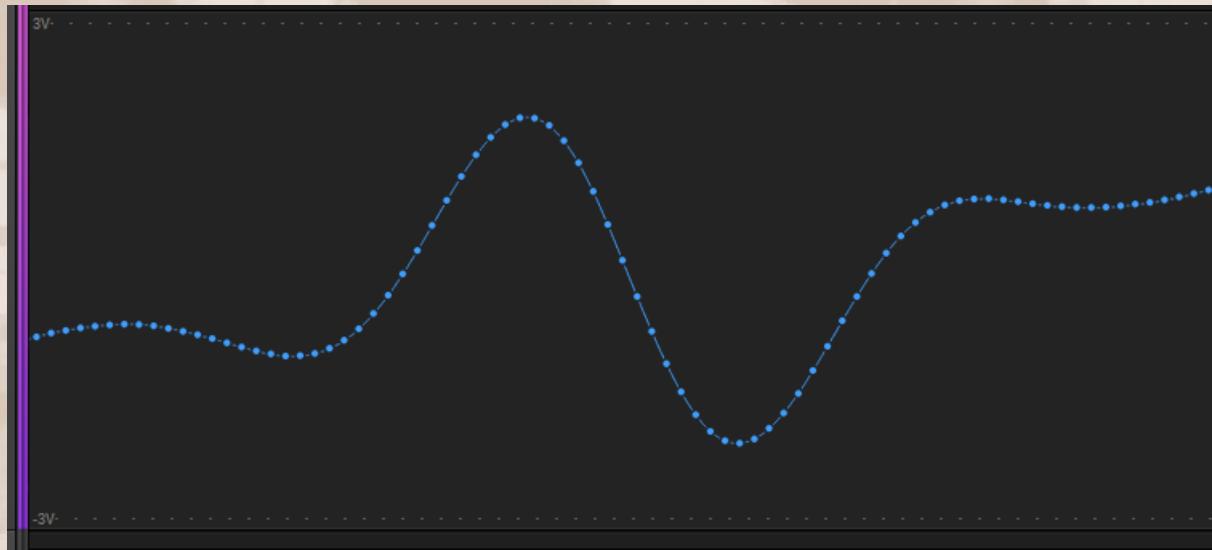
connected to a  
Qualstar 1052  
tape deck



6 to 16 tracks, 100 to 10,000 bits/inch



+3 volts



-3 volts

want 20 to 30 samples/bit

$$50 \text{ IPS} \times 800 \text{ BPI} \times 20 \text{ samples/bit} = 800\text{K samples/sec}$$

# the raw data

```
Time[s], 0 MSB, 1, 2, 3, 4, 5, 6, 7 LSB, 8P
Time [s],0 MSB-Analog,1-Analog,2-Analog,3-Analog,4-Analog,5-Analog,6-Analog,7 LSB-Analog,8P-Analog
4.507594559999999, 0.10085, -0.03400, -0.16784, -1.09570, 0.13550, 0.90018, 0.22597, 1.13374, -0.21236
4.507594880000000, -0.02258, 0.09941, -0.30091, -1.22020, 0.08454, 0.85893, 0.10245, 1.25193, -0.33846
4.507595200000000, -0.14601, 0.23281, -0.43398, -1.34989, 0.02849, 0.80220, -0.02622, 1.34956, -0.45448
4.507595520000000, -0.26944, 0.36109, -0.56193, -1.47439, -0.02247, 0.73515, -0.14974, 1.43179, -0.56545
4.507595840000000, -0.38773, 0.48937, -0.68477, -1.59370, -0.07342, 0.66295, -0.27326, 1.49859, -0.66633
4.507596159999999, -0.50087, 0.61252, -0.80248, -1.70264, -0.12438, 0.58044, -0.39677, 1.53970, -0.75713
4.507596479999999, -0.61401, 0.72540, -0.91508, -1.80120, -0.17534, 0.49793, -0.51000, 1.56539, -0.83279
4.507596800000000, -0.71687, 0.82802, -1.01232, -1.88420, -0.22629, 0.41541, -0.62323, 1.57567, -0.89332
4.507597120000000, -0.81459, 0.92038, -1.09933, -1.95163, -0.27215, 0.33290, -0.72616, 1.56025, -0.93872
4.507597440000000, -0.89687, 0.99735, -1.17098, -2.00351, -0.31292, 0.26070, -0.82395, 1.52942, -0.96898
4.507597759999999, -0.97402, 1.05379, -1.22728, -2.02945, -0.35368, 0.19881, -0.91144, 1.48831, -0.97907
4.507598079999999, -1.03573, 1.09997, -1.26823, -2.03982, -0.38935, 0.14208, -0.98349, 1.43179, -0.97403
4.507598400000000, -1.07688, 1.13076, -1.28870, -2.02426, -0.41483, 0.10598, -1.04011, 1.35984, -0.94881
4.507598720000000, -1.10259, 1.14102, -1.28870, -1.99313, -0.44031, 0.07504, -1.08643, 1.28790, -0.91350
4.507599040000000, -1.11288, 1.13076, -1.26823, -1.94126, -0.45050, 0.06473, -1.11216, 1.20568, -0.85801
4.507599359999999, -1.10259, 1.11023, -1.22216, -1.87382, -0.46069, 0.05957, -1.11730, 1.12346, -0.78739
4.507599679999999, -1.07173, 1.06405, -1.16075, -1.79601, -0.46069, 0.06988, -1.10701, 1.03096, -0.70669
4.507600000000000, -1.02030, 1.00761, -1.08398, -1.70264, -0.45050, 0.09051, -1.08128, 0.94360, -0.61085
4.507600320000000, -0.95345, 0.93577, -0.98673, -1.60407, -0.43012, 0.11630, -1.02981, 0.85110, -0.51501
4.507600640000000, -0.87116, 0.84855, -0.87925, -1.49514, -0.40464, 0.14724, -0.96291, 0.75861, -0.41917
4.507600960000000, -0.77344, 0.75619, -0.76666, -1.38101, -0.36897, 0.18334, -0.88571, 0.66611, -0.32333
4.507601279999999, -0.67059, 0.65356, -0.64382, -1.26689, -0.33330, 0.21428, -0.78792, 0.57361, -0.23253
```

CSV text file; can be over 50 GB

# compressed format

- binary “.tbin” file\*
- 16-bit scaled integers
- self-describing header
- utility program to convert

*10 to 1 space compression*

*2x faster processing*

\*defined with CHM software curator Al Kossow



# 7000 lines of C

```
void ww_assemble_data(void) { // assemble the array of 2-bit characters into bytes
    struct results_t *result = &block.results[block.parmset]; // where we put the results of this decoding
    uint16_t temp_data[MAXBLOCK + 1];
    int outndx = 0, nibble_counter = 0;
    uint16_t accum;
    // special hack: if there is one more clock than a multiple of 8, then assume the first clock is noise
    // and discard the first two bits we derived from it.
    if (ww.datacount % 8 == 1 && ww.datacount >= 9) {
        for (int ndx = 0; ndx < ww.datacount - 1; ++ndx)
            data[ndx] = data[ndx + 1];
        --ww.datacount;
        result->ww_leading_clock = 1; }
    if (reverse_tape) { // assemble going backwards, into a temp array
        for (int inndx = ww.datacount - 1; inndx >= 0; --inndx) {
            //accum = (accum >> 2) | ((data[inndx] & 0x03) << 6); // shift in 2 more bits, least significant first
            accum = (accum << 2) | (data[inndx] & 0x03); // shift in 2 more bits, most significant first
            if (++nibble_counter % 4 == 0) { // dump a full byte
                temp_data[outndx++] = (accum & 0xff) << 1; } } // create dummy parity bit on the right side
        for (int inndx = 0; inndx < outndx; ++inndx) // copy the temp array into the final data result
            data[inndx] = temp_data[inndx]; }
    else // assemble in place going forward, since we're making it smaller by 4x
        for (int inndx = 0; inndx < ww.datacount; ++inndx) {
            accum = (accum << 2) | (data[inndx] & 0x03); // shift in 2 more bits, most significant first
            if (++nibble_counter % 4 == 0) { // dump a full byte
                data[outndx++] = (accum & 0xff) << 1; } } // create dummy parity bit on the right side
    result->minbits = result->maxbits = outndx;
    if (ww.datacount % 8 != 0) { // should be a multiple of 16 bits, or 8 2-bit characters
        ++result->ww_bad_length;
        if (!doing_deskew && ww.datacount > 8 ) rlog(" *** the datacount for the next block is %d 2-bit characters, "
            "which is %d more than a multiple of 8\n",
            ww.datacount, ww.datacount % 8); }
    float target_bitspace = 1 / (bpi*ips);
    if (fabs(ww.clkavg.t_bitspaceavg - target_bitspace) / target_bitspace > WW_MAX_CLK_VARIATION) ++result->ww_speed_err; }
```

[github.com/LenShustek/readtape](https://github.com/LenShustek/readtape)

# No GUI, but lots of options

```
-ntrks=n      set the number of tracks
-order=       set input data order for tracks 0..ntrks-2,P, where 0=MSB
              default: 01234567P for 9 trk, 012345P for 7 trk
              (for Whirlwind: a combination of C L M c l m and x's)

-pe          PE (phase encoding)
-nrzi        NRZI (non return to zero inverted)
-gcr         GCR (group coded recording)
-whirlwind I  Whirlwind I 6-track 2-bit-per-character
-ips=n       speed in inches/sec (default: 50, except 25 for GCR)
-bpi=n       density in bits/inch (default: autodetect)
-zeros       base decoding on zero crossings instead of peaks
-even        expect even parity instead of odd (for 7-track NRZI BCD tapes)
-invert      invert the data so positive peaks are negative and vice versa
-fluxdir=d   flux direction is 'pos', 'neg', or 'auto' for each block
-reverse     reverse bits in a word and words in a block (Whirlwind only)
-skip=n      skip the first n samples
-blklimit=n  stop after n blocks
-subsample=n use only every nth data sample
-showibg=n   report on interblock gaps greater than n milliseconds
-tap         create one SIMH .tap file from all the data
-deskew      do NRZI track deskewing based on the beginning data
-skew=n,n    use this skew, in #samples for each track, rather than deducing it
-correct     do error correction, where feasible
-addparity   include the parity bit as the highest bit in the data (for ntrks<9)
-tbin        only look for a .tbin input file, not .csv first
-nolog       don't create a log file
-nolabels    don't try to decode IBM standard tape labels
-textfile    create an interpreted .<options>.txt file from the data
              numeric options: -hex -octal (bytes) -octal2 (16-bit words)
              character options: -ASCII -EBCDIC -BCD -sixbit -B5500 -SDS -SDSM -flexo
              characters per line: -linesize=nn
              space every n bytes of data: -dataspace=n
              make LF or CR start a new line: -linefeed

-outf=bbb    use bbb as the <basefilename> for output files
-outp=ppp    otherwise use ppp as an optional prepended path for output files
-sumt=sss    append a text summary of results to text file sss
-sumc=ccc    append a CSV summary of results to text file ccc
-m           try multiple ways to decode a block
-nm          don't try multiple ways to decode a block
-v[n]        verbose mode [level n, default is 1]
-q           quiet mode (only say "ok" or "bad")
-f           take a file list from <basefilename>.txt
```

# decode each block with different parameters until error-free

reading parmsets from "PE.parms"

	clk_window,	clk_alpha,	agc_window,	agc_alpha,	min_peak,	clk_factor,	pulse_adj,	pkww_bitfrac
{	0,	0.200,	5,	0.000,	0.000,	1.500,	0.400,	0.700, },
{	0,	0.200,	5,	0.000,	0.100,	1.500,	0.400,	0.700, },
{	3,	0.000,	5,	0.000,	0.000,	1.400,	0.000,	0.700, },
{	3,	0.000,	5,	0.000,	0.000,	1.400,	0.200,	0.700, },
{	5,	0.000,	5,	0.000,	0.000,	1.400,	0.000,	0.700, },
{	5,	0.000,	5,	0.000,	0.000,	1.500,	0.200,	0.700, },
{	5,	0.000,	5,	0.000,	0.000,	1.400,	0.400,	0.700, },
{	3,	0.000,	5,	0.000,	0.000,	1.400,	0.200,	0.700, }





slow!

# dataflow

.logicdata  
file

.csv file

csvtbin

.tbin

readtape

speed on a laptop:  
about real time

.bin

.tap

Supnik SIMH  
format


.log

.txt

.csv

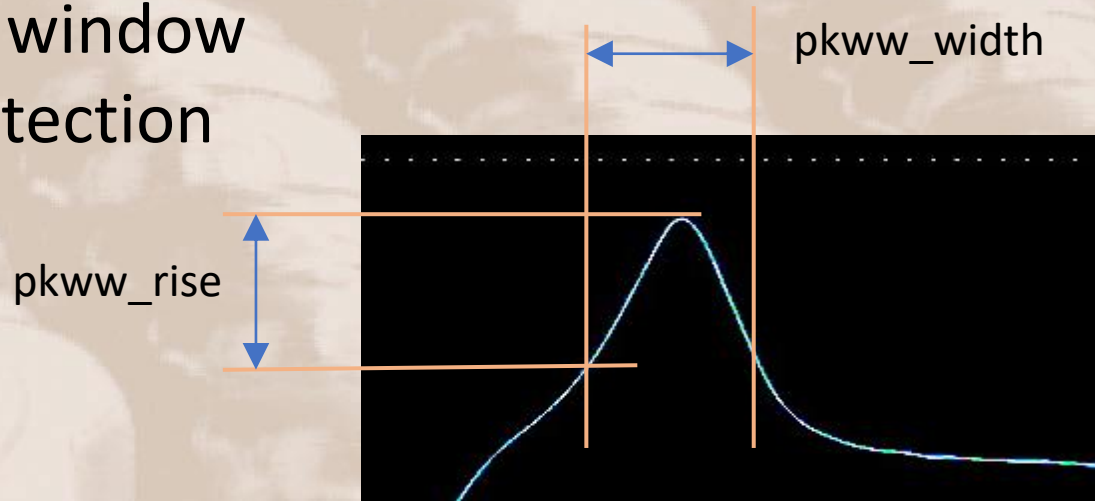
for debugging

# debugging

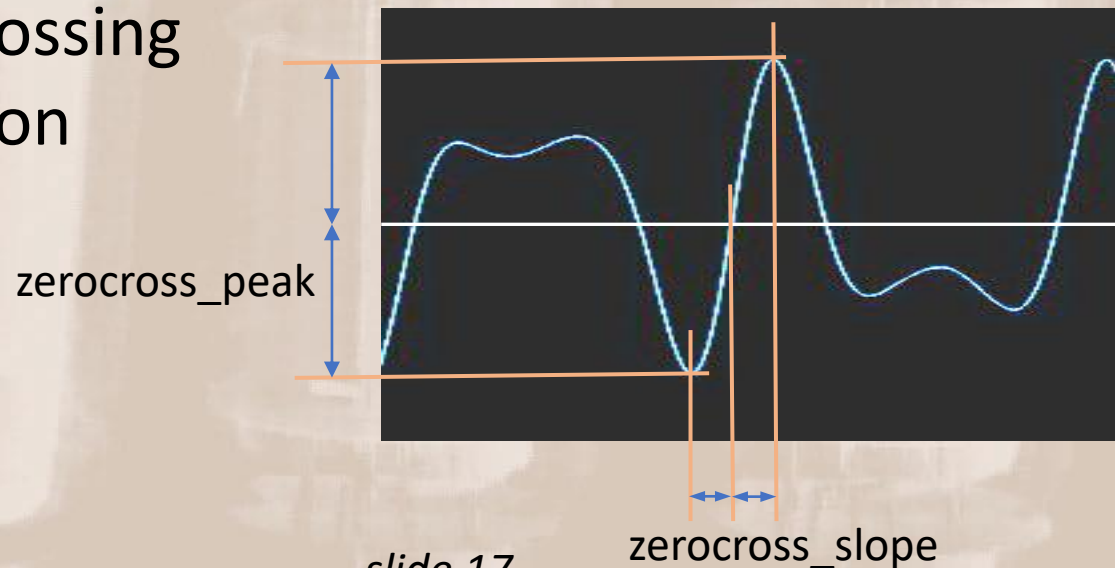


# detecting transitions

option 1: moving window  
peak detection



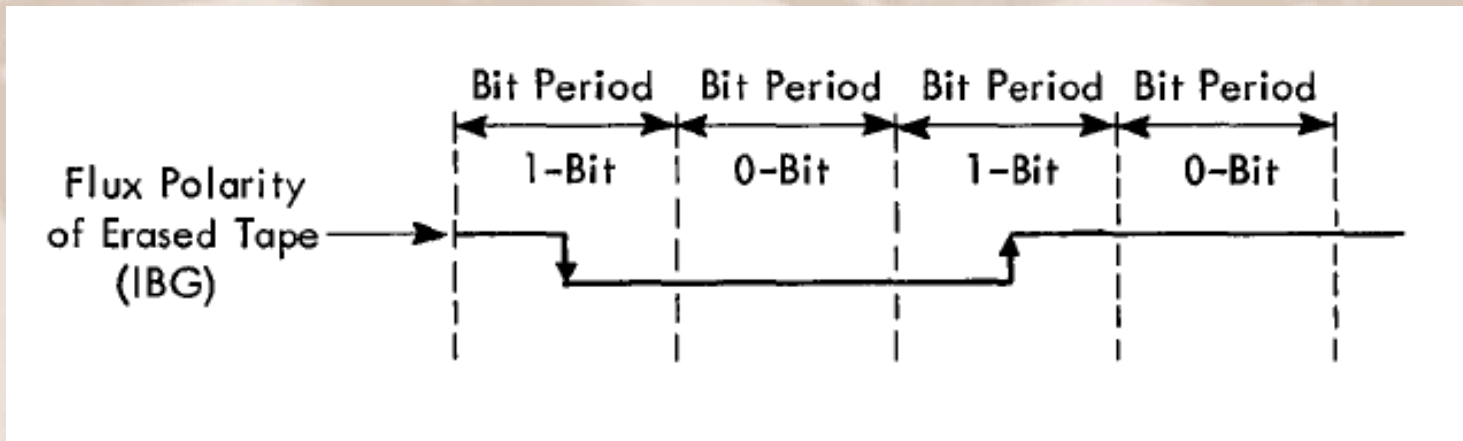
option 2: zero-crossing  
detection





# decoding the bits

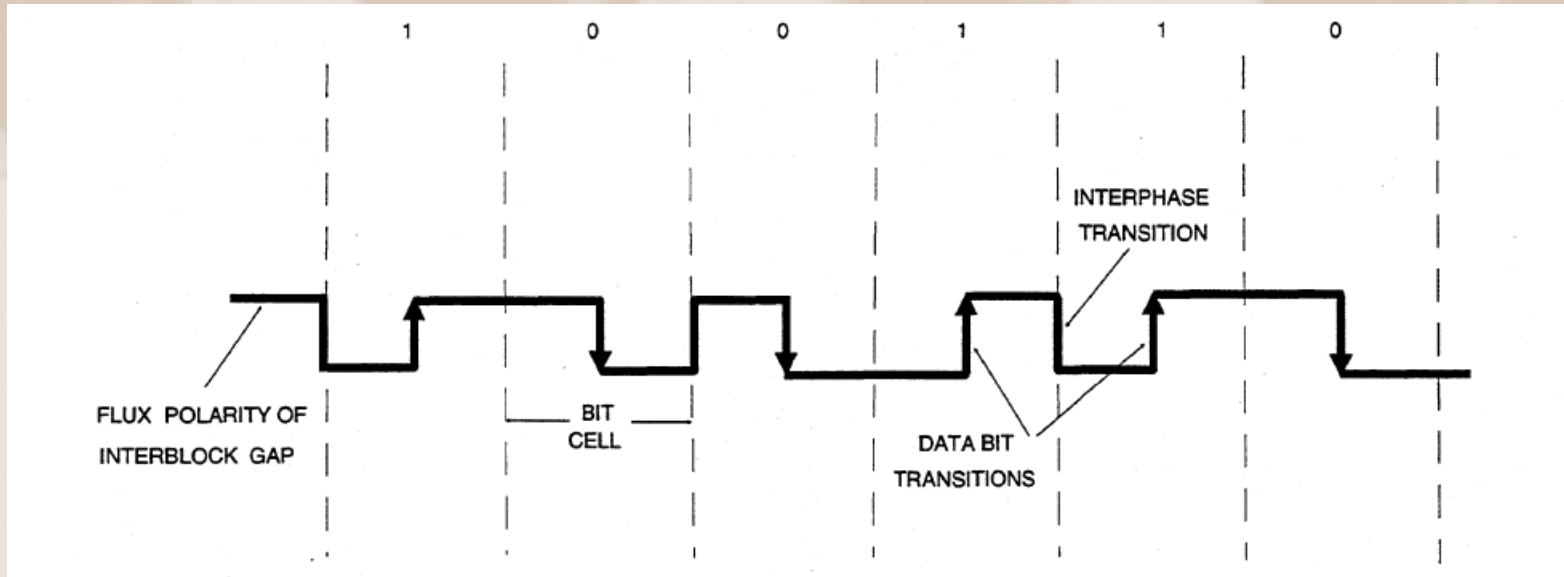
## NRZI: Non return to Zero Inverted



- used for 7-track, 9-track 800 BPI (and 6250 BPI GCR), and Whirlwind tapes
- NOT self-clocking – requires odd parity or group recoding
- often: odd parity, LRC and CRC block check characters
- sometimes: preamble of 40 zeroes followed by 1
- sometimes: filemark is 64-256 flux reversals on tracks 2, 5, 8 only

# decoding the bits

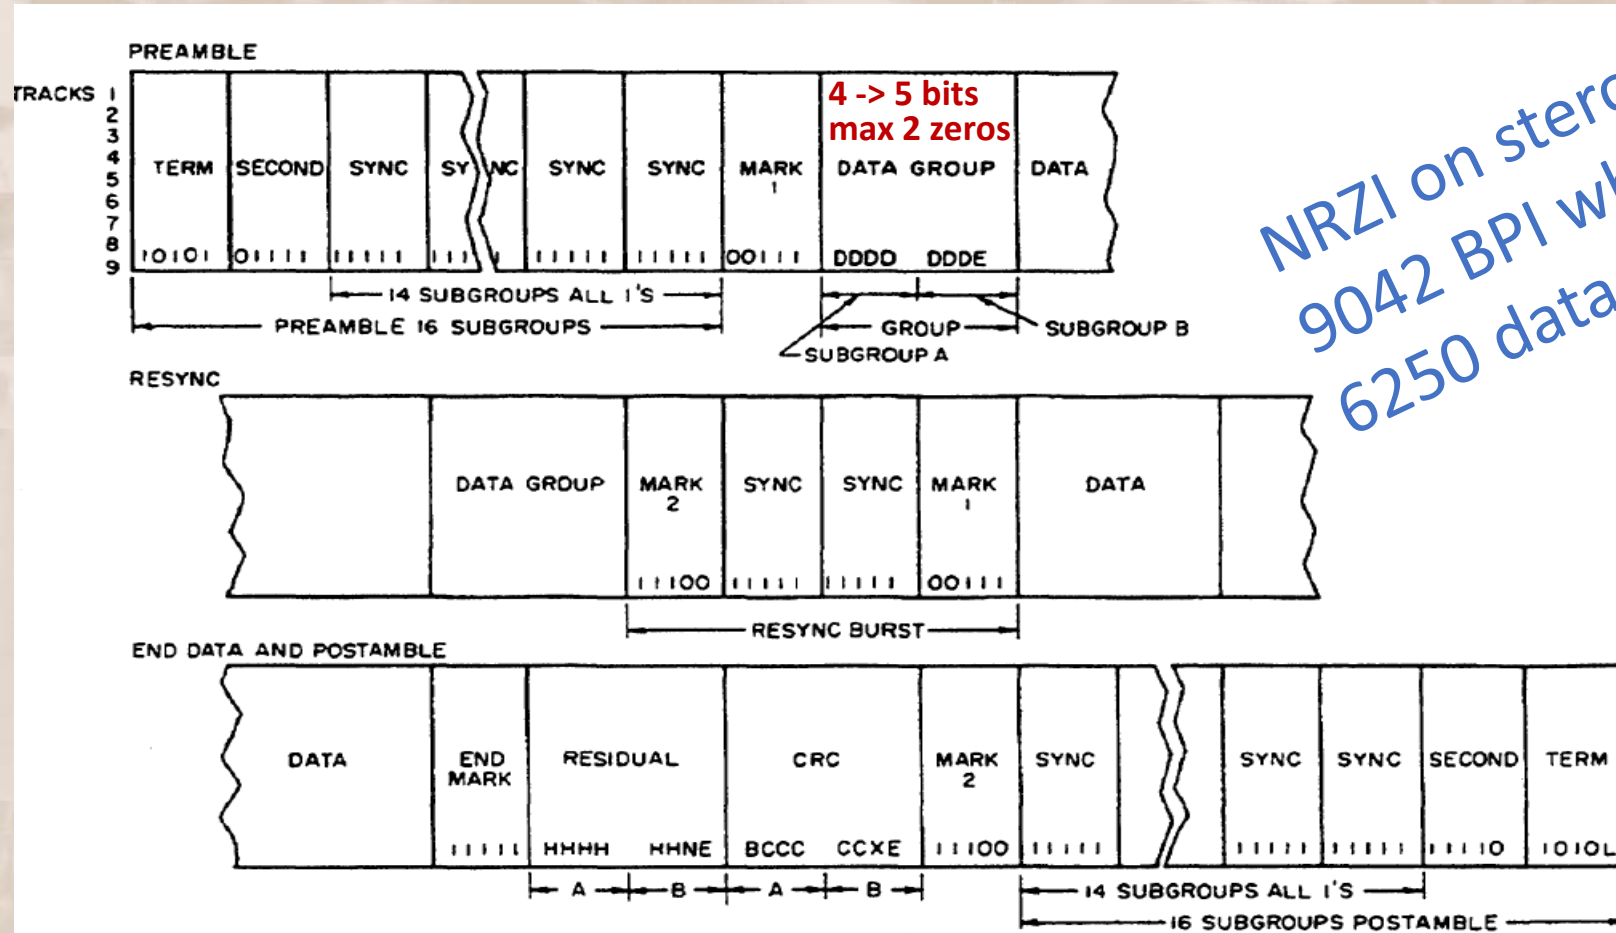
## PE: Phase Encoding



- used for 9-track 1600 BPI drives
- self-clocking
- preamble of 40 zeroes followed by 1
- odd parity, no longitudinal block checks
- filemark: 64-256 flux reversals on tracks 2, 5, 8 only


# decoding the bits

## GCR: Group Coded Recording






# fixing head skew




pre-read 1000 transitions  
to determine head skew,  
then delay track data


before deskewing



after deskewing



# fixing amplitude changes



AGC (Automatic Gain Control) based on:

- average height of last N peaks, or
- exponential weighted average of previous peaks

$$\text{gain} = \alpha \times (\text{height}/\text{avgheight}) + (1 - \alpha) \times \text{gain}$$

# fixing speed changes

Initial bit clock estimate is:

$$1 / ((\text{bits/inch}) * \text{inches/sec})$$

Update it after each bit to:

- same
- average spacing of last N bits, or
- exponential weighted average of previous bit spacings

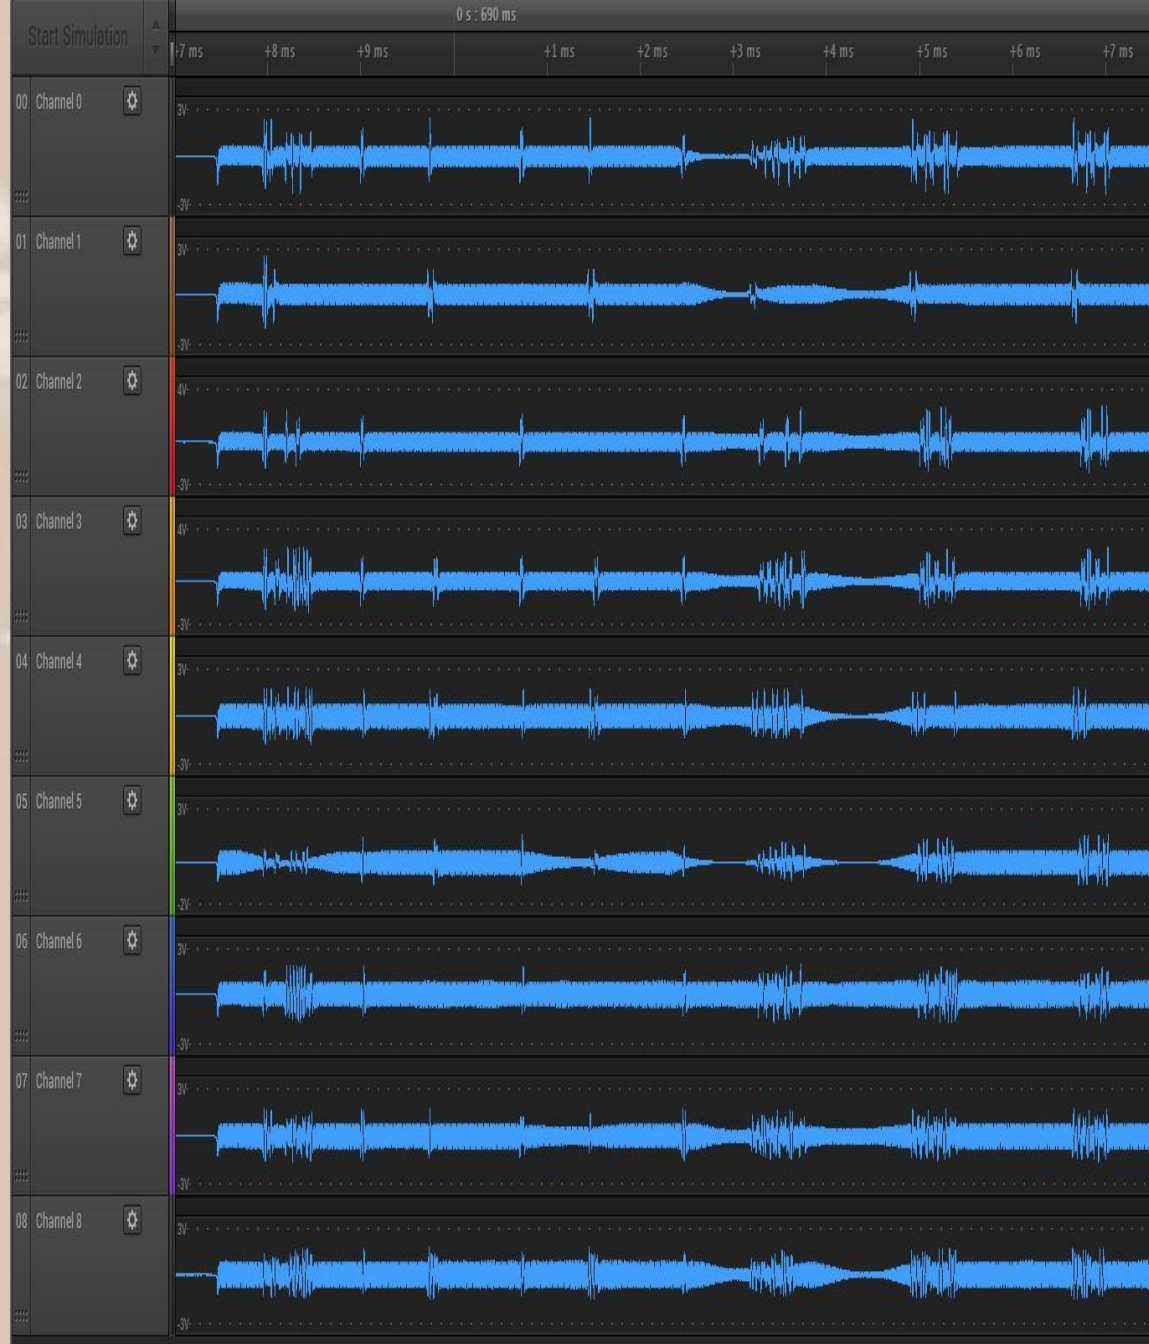
$$\text{clock} = \alpha \times \text{last\_bit\_spacing} + (1 - \alpha) \times \text{clock}$$





# fixing errors

- deduce the bits from parity, or
- correct the bits from ECC codes, or
- during dropout, fake “same as before” bits



# result: PLAGO 100% recovered

```
reading file "plago_whole_02.csv" on Wed Mar 28 21:22:57 2018
*** tape label VOL1: PLAGOS, owner P.I.B.
*** tape label HDR1: DOCUMENT , serno PLAGOS, created 70181
    volume 0001, dataset 0001
creating file "plago_whole_02\000-DOCUMENT.bin"
*** tape label HDR2: RECFM=FB, BLKSIZE=00800, LRECL=00080
    job: LENNY /STEP1
*** tapemark
wrote block 1, 800 bytes, 1 tries, parms 0, max AGC 1.35, 0 parity errs, CRC ok , LRC ok , avg speed 50.10 IPS, at time 3.5078912
wrote block 2, 800 bytes, 1 tries, parms 0, max AGC 1.37, 0 parity errs, CRC ok , LRC ok , avg speed 49.69 IPS, at time 3.5398886
wrote block 3, 800 bytes, 1 tries, parms 0, max AGC 1.25, 0 parity errs, CRC ok , LRC ok , avg speed 49.51 IPS, at time 3.5722214
wrote block 4, 800 bytes, 1 tries, parms 0, max AGC 1.32, 0 parity errs, CRC ok , LRC ok , avg speed 50.12 IPS, at time 3.6038989...
```

```
***** 00081400
* 00081500
***** FSB    FETCH SUBSCRIPTED REFERENCE ***** 00081600
* 00081700
***** 00081800
    SPACE 00081900
EFSB    LH    TR4,0 (LOCNTR)    OFFSET TO ID ENTRY    00082000
        A    TR4,IDBASE    ABS ADDR OF ID ENTRY    00082100
        LA    TR13,ENEXT    SET RETURN ADDRESS    00082200
        LH    TR5,2 (LOCNTR)    GET NO. OF SUBSCRIPTS    00082300
        LA    LOCNTR,4 (LOCNTR)    00082400
EFSB1    DS    0H    ENTER HERE FROM GTD RTN, TR5=NO. OF SUBS, TR4=ID    00082500
        LH    TR3,10 (TR4)    OFFSET OF DV IN VSTK    00082600
        L    TR2,VBLKTBL    NO - GET ADDRESS OF BLOCK TABLE    00082700
        A    TR3,0 (TR2,EOPERND)    ADD PROPER VSTK BASE    00082800
        TM    8 (TR4) ,PARAM    IS THIS A PARAMTTER    00082900
        BZ    EFSB9    NO...    00083000
        A    TR3,0 (,TR3)    YES-GET PTR TO ARGUMENT    00083100
        BNP    ERR22    BR IF ARG WAS NOT PASSED    00083200
EFSB9    L    TR2,0 (,TR3)    LOAD VIRTUAL ORIGIN    00083300
        LA    TR0,0 (TR2)    * INTO ACCUMULATOR (TR0)    00083400
        BCTR    TR5,0    SET TR5 = 4 * (NO. OF SUBS - 1)    00083500
        SLA    TR5,2    *    00083600
        SR    DSTK,TR5    PURGE SUBS, BUT LEAV A SLOT    00083700
        LR    TR2,DSTK    POINT TO FIRST SUBSCRIPT    00083800
        ST    CVSTK,SAV1    SAVE CURRENT VSTK POINTER    00083900
```

# other tapes decoded


- Englebart NLS (SDS-940)
- SRI PDP-10
- Microdata Reality OS
- Lawrence Berkeley Lab
- SEL realtime OS
- Project Genie timesharing (UCB)
- BBN tapes
- etc.

## why do this?

- recover lost knowledge about early computers and software
- show techniques for digital archaeology
- provide authentic programs and data for restorations and simulations




COMPUTER  
STRUCTURES:  
READINGS  
AND  
EXAMPLES  
C GORDON BELL & ALLEN NEWELL



# case study: Whirlwind

(the inspiration for CHM?)

Gordon Bell:  
Somebody should  
preserve the  
historic computers!

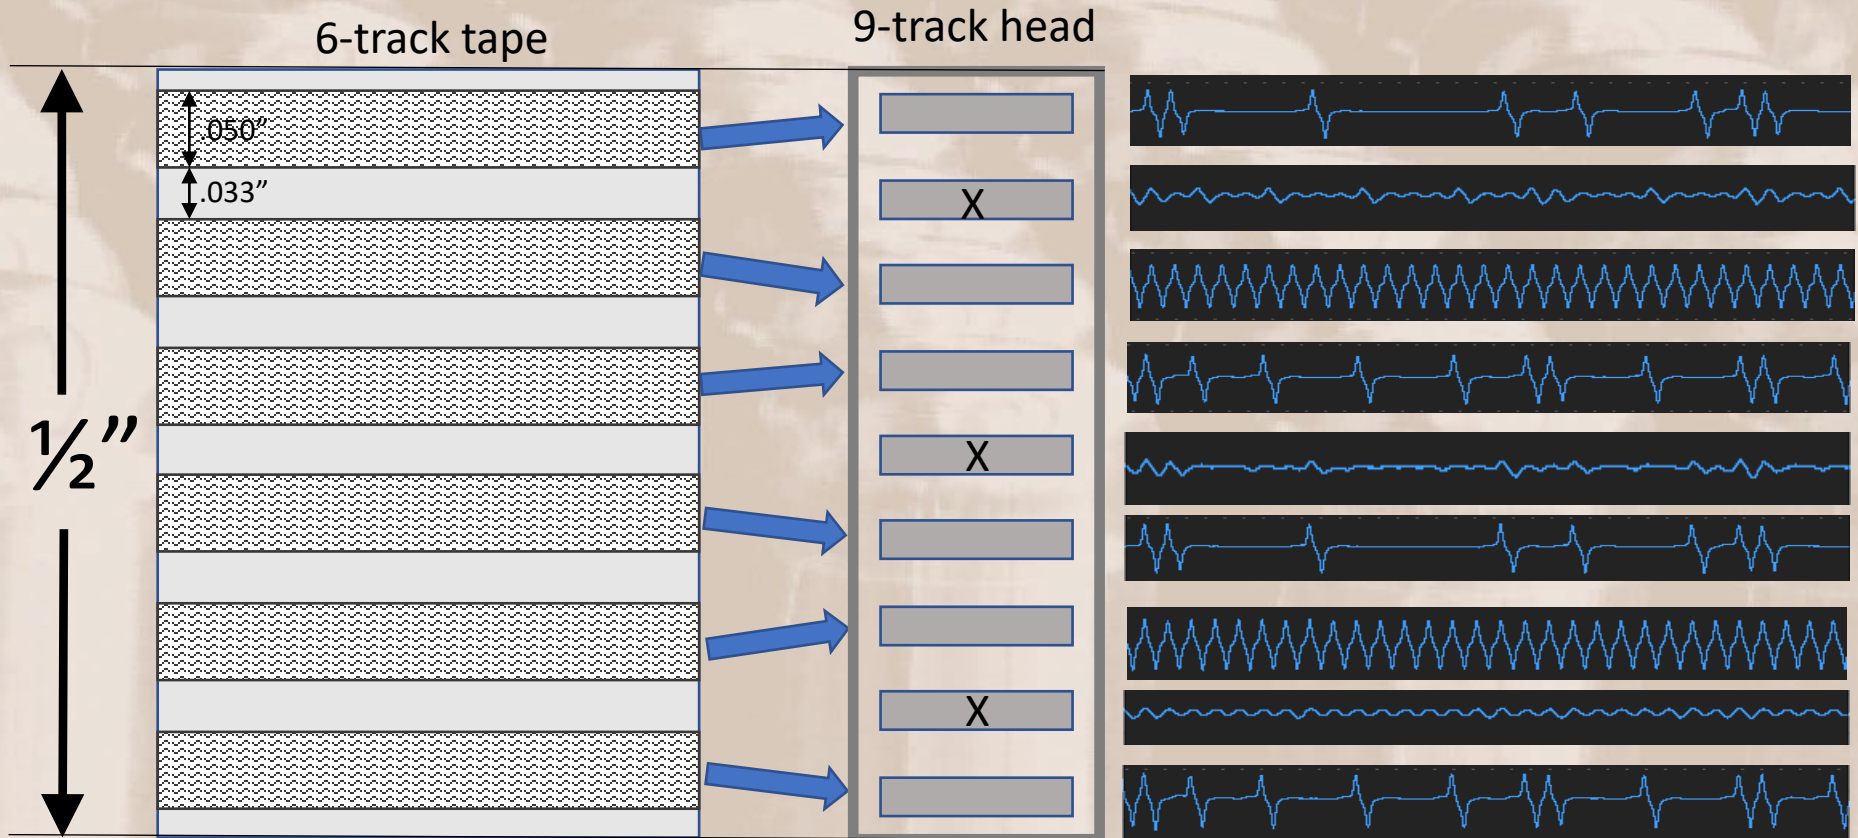


# 40+ years later, the 150 6-track Whirlwind tapes were still unread





# 6 track head -- NOT



the head misalignment is always only 0.014"  
from the center of the 0.050" wide track



# decoding Whirlwind bits



MSB

clock

LSB

LSB (copy)

clock (copy)

MSB (copy)

- basically NRZI
- 2 bits/character assembled into 16-bit words
- each track is duplicated
- no parity or other checks
- the flux transition polarity varies from tape to tape
- tape mark: LSB bits without a clock

# dataflow



## results

- decoded 61 tapes so far
  - about half the inventory
- 161,955 blocks
  - 0.13% had errors
- 10,274,572 bytes
- data and source code
  - very little text ☹️

*stay tuned for Guy's talk about using the bits*

# questions?

[shustek@computerhistory.org](mailto:shustek@computerhistory.org)